

Praful_file1_hw2

February 15, 2024

Spam Detection HW

Read complete instructions before starting the HW

1 Installing/Importing Modules

```
[36]: !pip install -U spacy -q
```

```
[37]: !python -m spacy download en_core_web_sm
```

Collecting en-core-web-sm==3.7.1

Downloading https://github.com/explosion/spacy-models/releases/download/en_core_web_sm-3.7.1/en_core_web_sm-3.7.1-py3-none-any.whl (12.8 MB)

12.8/12.8 MB

23.6 MB/s eta 0:00:00

Requirement already satisfied: spacy<3.8.0,>=3.7.2 in /usr/local/lib/python3.10/dist-packages (from en-core-web-sm==3.7.1) (3.7.2)

Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.0.12)

Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.0.5)

Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.0.10)

Requirement already satisfied: cymem<2.1.0,>=2.0.2 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.0.8)

Requirement already satisfied: preshed<3.1.0,>=3.0.2 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.0.9)

Requirement already satisfied: thinc<8.3.0,>=8.1.8 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (8.2.3)

Requirement already satisfied: wasabi<1.2.0,>=0.9.1 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-

sm==3.7.1) (1.1.2)

Requirement already satisfied: srsly<3.0.0,>=2.4.3 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.4.8)

Requirement already satisfied: catalogue<2.1.0,>=2.0.6 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.0.10)

Requirement already satisfied: weasel<0.4.0,>=0.1.0 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.3.4)

Requirement already satisfied: typer<0.10.0,>=0.3.0 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.9.0)

Requirement already satisfied: smart-open<7.0.0,>=5.2.1 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (6.4.0)

Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (4.66.1)

Requirement already satisfied: requests<3.0.0,>=2.13.0 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.31.0)

Requirement already satisfied: pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.6.1)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages
 (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.1.3)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages
 (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (67.7.2)

Requirement already satisfied: packaging>=20.0 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (23.2)

Requirement already satisfied: langcodes<4.0.0,>=3.2.0 in
 /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.3.0)

Requirement already satisfied: numpy>=1.19.0 in /usr/local/lib/python3.10/dist-packages
 (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.25.2)

Requirement already satisfied: annotated-types>=0.4.0 in
 /usr/local/lib/python3.10/dist-packages (from
 pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.6.0)

Requirement already satisfied: pydantic-core==2.16.2 in
 /usr/local/lib/python3.10/dist-packages (from
 pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.16.2)

Requirement already satisfied: typing-extensions>=4.6.1 in
 /usr/local/lib/python3.10/dist-packages (from
 pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4->spacy<3.8.0,>=3.7.2->en-core-web-

```

sm==3.7.1) (4.9.0)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from
requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-
packages (from requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-web-
sm==3.7.1) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from
requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from
requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2024.2.2)
Requirement already satisfied: blis<0.8.0,>=0.7.8 in
/usr/local/lib/python3.10/dist-packages (from
thinc<8.3.0,>=8.1.8->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.7.11)
Requirement already satisfied: confection<1.0.0,>=0.0.1 in
/usr/local/lib/python3.10/dist-packages (from
thinc<8.3.0,>=8.1.8->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.1.4)
Requirement already satisfied: click<9.0.0,>=7.1.1 in
/usr/local/lib/python3.10/dist-packages (from
typer<0.10.0,>=0.3.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (8.1.7)
Requirement already satisfied: cloudpathlib<0.17.0,>=0.7.0 in
/usr/local/lib/python3.10/dist-packages (from
weasel<0.4.0,>=0.1.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.16.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.10/dist-packages (from jinja2->spacy<3.8.0,>=3.7.2->en-
core-web-sm==3.7.1) (2.1.5)
Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')

```

```
[38]: !pip install pyspellchecker
```

```

Requirement already satisfied: pyspellchecker in /usr/local/lib/python3.10/dist-
packages (0.8.1)

```

```

[39]: import spacy
from spacy.matcher import Matcher
from spacy.tokens import Token
import pandas as pd
import numpy as np
from nltk.stem.porter import PorterStemmer
import os
import sys
from pathlib import Path
from typing import List
from sklearn.pipeline import Pipeline, FeatureUnion

```

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from collections import Counter
from xgboost import XGBClassifier
from sklearn.metrics import fbeta_score, make_scorer
from sklearn.model_selection import RandomizedSearchCV
from sklearn.base import BaseEstimator, TransformerMixin
from bs4 import BeautifulSoup
import re
from spellchecker import SpellChecker
import warnings
warnings.filterwarnings('ignore')

```

```

[40]: from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

[41]: base_folder = Path('/content/drive/MyDrive/NLP_HW/')
data_folder = base_folder/'HW_2'

```

2 Q1: Load the dataset (1 Point)

- For this Hw you will use spam dataset from kaggle which can be found from [this](#) link. You can download this data and either upload it in google drive or in colab workspace. Load the data in pandas dataframe.
- There are only two useful columns. These columns are related to (1) label (ham and spam) and the (2) text of email.
- Rename columns as label and message
- Find the % ham and spam in the data.

```

[42]: file = data_folder/'spam.csv'
df = pd.read_csv(file, encoding = 'ISO-8859-1')
df = df[['v1', 'v2']]
df = df.rename(columns = {'v1': 'label', 'v2': 'message'})

```

```

[43]: df.head()

```

```

[43]:   label      message
0   ham  Go until jurong point, crazy.. Available only ...
1   ham                      Ok lar... Joking wif u oni...
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...
3   ham  U dun say so early hor... U c already then say...
4   ham  Nah I don't think he goes to usf, he lives aro...

```

```
[44]: percent_ham = df['label'].value_counts()['ham']/len(df['label'])*100
      percent_spam = df['label'].value_counts()['spam']/len(df['label'])*100
      print("Percentage of ham:", percent_ham)
      print("Percentage of spam:", percent_spam)
```

Percentage of ham: 86.59368269921033

Percentage of spam: 13.406317300789663

3 Q2 : Provide the metric for evaluating model (1 Point)

As you will notice, the data is highly imbalanced (most messages are labelled as ham and only few are labelled as spam). Always predicting ham will give us very good accuracy (close to 90%). So you need to choose a different metric.

Task: Provide the metric you will choose to evaluate your model. Explain why this is an appropriate metric for this case.

3.0.1 Model Evaluation Metric

Since the data is highly biased, accuracy is not a good performance metric for evaluating models of this dataset since even a dummy model which classifies all observations as ham would generate an accuracy of approximately 86.59%.

F1 Score and F2 Score are better evaluation metrics for such an imbalanced dataset.

The F1 Score is given by $(2 * (\text{Precision} * \text{Recall})) / (\text{Precision} + \text{Recall})$ where Precision is $(\text{True Positives}) / (\text{True Positives} + \text{False Negatives})$ and Recall is $(\text{True Positives}) / (\text{True Positives} + \text{False Negatives})$.

Precision limits the number of False Positives and Recall identifies all positive samples. A harmonic mean of the two is thus a good measure in our case. It **penalizes** naive/dummy models.

In F1 Score, both Precision and Recall are given an equal weightage but in our case it is more important to ensure that an actual spam message is not missed out. Therefore, we need to provide more weightage to Recall since it identifies all positive samples. **F2 score** gives more weightage to Recall and is thus a good metric for our classification problem.

F2 score is given by $(\text{True Positives}) / ((\text{True Positives} + (0.2 * \text{False Positives}) + (0.8 * \text{False Negatives})))$

By having more weightage to False Negatives in the denominator, we are emphasizing more on the smaller proportion of our data (spam messages).

4 Q3 : Classification Pipelines (18 Points)

In the previous lectures you learned Data processing, Featurization such as CountVectorizer, TfidfVectorizer, and also Feature Engineering. * You will now use following methods to create features which you can use in your model.

1. Sparse Embeddings (TF-IDF) (6 Points)
2. Feature Engineering (see examples below) (6 Points)

3. Sparse Embeddings (TF-IDF) + Feature Engineering (6 Points)

Approach:

****Use a smaller subset of dataset (e.g. 5-10 %) to evaluate the three pipelines . Based on your analysis (e.g. model score, learning curves) , choose one pipeline from the three. Provide your rationale for choosing the pipeline. Train only the final pipeline on randomly selected larger subset (e.g. 40%) of the data.**

Requirements:

1. You can use any ML model (Logistic Regression, XgBoost) for the classification. You will need to tune the **model for imbalanced dataset** (The link on XGBoost tutorial for imbalanced data: <https://machinelearningmastery.com/xgboost-for-imbalanced-classification/>).
2. For feature engineering, you can choose from the examples below. You do not have to use all of them. You can add other features as well. Think about what features can distinguish a spam from a regular email. Some examples :

Count of following (Words, characters, digits, exclamation marks, numbers, Nouns, ProperNouns, AUX, VERBS, Adjectives, named entities, spelling mistakes (see the link on how to get spelling mistakes <https://pypi.org/project/pyspellchecker/>).

3. For Sparse embeddings you will use **tfidf vectorization**. You need to choose appropriate parameters e.g. min_df, max_df, max_features, n-grams etc.).
4. Think carefully about the pre-processing you will do.

Tip: **Using GridSearch for hyperparameter tuning might take a lot of time. Try using RandomizedSearch.** You can also explore faster implementation of Gridsearch and Randomized-Search in sklearn:

1. [Halving Grid Search](#)
2. [HalvingRandomSearchCV](#)

4.0.1 Taking a subset of the dataset and splitting it to create train and test datasets

```
[45]: df_small = df.sample(frac = 0.1, random_state = 0)
      df_small
```

```
[45]:      label      message
4456  ham  Aight should I just plan to come up later toni...
690   ham                        Was the farm open?
944   ham  I sent my scores to sophas and i had to do sec...
3768  ham  Was gr8 to see that message. So when r u leavi...
1189  ham  In that case I guess I'll see you at campus lodge
...   ...
4587  ham  I wanted to wish you a Happy New Year and I wa...
152   ham                        Ok... Ur typical reply...
2357  ham  Okay same with me. Well thanks for the clarifi...
1559  ham  Single line with a big meaning::::: \Miss anyt...
```

```
3544    ham                                Thank You meet you monday
```

```
[557 rows x 2 columns]
```

```
[46]: df_small['label'] = df_small['label'].map({'spam':1, 'ham':0}).astype(int)
```

```
[47]: X = df_small['message'].values
      y = df_small['label'].values
      X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

We shall use F2 Score as the evaluation metric for all our pipelines

```
[48]: f2score = make_scorer(fbeta_score, beta=2)
```

4.0.2 Sparse Embeddings

```
[49]: counter = Counter(y)
      estimate = counter[0] / counter[1]
```

```
[50]: pipe1 = Pipeline([('sparse_embed', TfidfVectorizer()), ('classifier',
      ↪XGBClassifier(scale_pos_weight=estimate))])
      param_grid = {'sparse_embed_max_df': [0.15, 0.2, 0.5, 0.75, 1],
      ↪'sparse_embed_max_features': [None, 100, 200, 400, 500, 800]}
```

```
[51]: random_search1 = RandomizedSearchCV(estimator=pipe1,
      ↪param_distributions=param_grid, cv = 5, scoring= f2score)
      random_search1.fit(X_train, y_train)
```

```
[51]: RandomizedSearchCV(cv=5,
                        estimator=Pipeline(steps=[('sparse_embed',
                                                    TfidfVectorizer()),
                                                    ('classifier',
                                                     XGBClassifier(base_score=None,
                                                                    booster=None,
                                                                    callbacks=None,

colsample_bylevel=None,
colsample_bynode=None,
colsample_bytree=None,

device=None,

early_stopping_rounds=None,
enable_categorical=False,

eval_metric=None,
feature_types=None,
gamma=None,
grow_policy=Non...
max_depth=None,
max_leaves=None,
```

```

min_child_weight=None,
monotone_constraints=None,
num_parallel_tree=None,
...))),
param_distributions={'sparse_embed__max_df': [0.15, 0.2, 0.5,
                                              0.75, 1],
                    'sparse_embed__max_features': [None,
                                                    100, 200,
                                                    400, 500,
                                                    800]},
scoring=make_scorer(fbeta_score, beta=2))

```

```

[52]: print("Best Parameters: ", random_search1.best_params_)
      print("Best F2 score:", random_search1.best_score_)
      print("Best Estimator: ", random_search1.best_estimator_)

```

```

Best Parameters:  {'sparse_embed__max_features': 500, 'sparse_embed__max_df':
0.15}
Best F2 score: 0.7359480145431272
Best Estimator: Pipeline(steps=[('sparse_embed',
TfidfVectorizer(max_df=0.15, max_features=500)),
('classifier',
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None, gamma=None, grow_policy=None,
              importance_type=None,
              interaction_constraints=None, learning_rate=None,
              max_bin=None, max_cat_threshold=None,
              max_cat_to_onehot=None, max_delta_step=None,
              max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None, multi_strategy=None,
              n_estimators=None, n_jobs=None,
              num_parallel_tree=None, random_state=None,
...))])

```

```

[53]: print("The train score is", random_search1.score(X_train, y_train))
      print("The test score is", random_search1.score(X_test, y_test))

```

The train score is 1.0

The test score is 0.7216494845360824

4.0.3 Feature Engineering

We will use the featurizer class that was provided in the lecture along with few additions to count exclamations and misspelled words. We will also include the definition for the Custom Preprocessor class since this is used in the featurizer class.

Spacy Preprocessor

```
[54]: class SpacyPreprocessor(BaseEstimator, TransformerMixin):  
  
    """  
    A text preprocessor that utilizes spaCy for efficient and flexible NLP.▯  
    ↪Designed as a part of a scikit-learn  
    pipeline, it provides a wide range of text cleaning and preprocessing▯  
    ↪functionalities.  
  
    Attributes:  
        model (str): The spaCy language model to be used for tokenization and▯  
    ↪other NLP tasks.  
        batch_size (int): The number of documents to process at once during▯  
    ↪spaCy's pipeline processing.  
        lemmatize (bool): If True, lemmatize tokens.  
        lower (bool): If True, convert all characters to lowercase.  
        remove_stop (bool): If True, remove stopwords.  
        remove_punct (bool): If True, remove punctuation.  
        remove_email (bool): If True, remove email addresses.  
        remove_url (bool): If True, remove URLs.  
        remove_num (bool): If True, remove numbers.  
        stemming (bool): If True, apply stemming to tokens (mutually exclusive▯  
    ↪with lemmatization).  
        add_user_mention_prefix (bool): If True, add '@' as a separate token▯  
    ↪(useful for user mentions in social  
        media data).  
        remove_hashtag_prefix (bool): If True, do not separate '#' from the▯  
    ↪following text.  
        basic_clean_only (bool): If True, perform only basic cleaning (HTML▯  
    ↪tags removal, line breaks, etc.)  
        and ignore other preprocessing steps.  
  
    Methods:  
        basic_clean(text: str) -> str:  
            Performs basic cleaning of the text such as removing HTML tags and▯  
    ↪excessive whitespace.  
  
        spacy_preprocessor(texts: list) -> list:
```

Processes a list of texts through the spaCy pipeline with specified,
↳ preprocessing options.

`fit(X, y=None) -> 'SpacyPreprocessor':`

Fits the preprocessor to the data. This is a dummy method for,
↳ scikit-learn compatibility and does not
change the state of the object.

`transform(X, y=None) -> list:`

Transforms the provided data using the defined preprocessing,
↳ pipeline. Performs basic cleaning,
and if `basic_clean_only` is False, it applies advanced spaCy,
↳ preprocessing steps.

Raises:

`ValueError`: If both 'lemmatize' and 'stemming' are set to True.

`ValueError`: If 'basic_clean_only' is True but other processing options,
↳ are also set to True.

`TypeError`: If the input X is not a list or a numpy array.

"""

`def __init__(self, model, *, batch_size = 64, lemmatize=True, lower=True,
↳ remove_stop=True,`

`remove_punct=True, remove_email=True, remove_url=True,
↳ remove_num=False, stemming = False,
add_user_mention_prefix=True, remove_hashtag_prefix=False,
↳ basic_clean_only=False):`

`self.model = model`

`self.batch_size = batch_size`

`self.remove_stop = remove_stop`

`self.remove_punct = remove_punct`

`self.remove_num = remove_num`

`self.remove_url = remove_url`

`self.remove_email = remove_email`

`self.lower = lower`

`self.add_user_mention_prefix = add_user_mention_prefix`

`self.remove_hashtag_prefix = remove_hashtag_prefix`

`self.basic_clean_only = basic_clean_only`

`if lemmatize and stemming:`

`raise ValueError("Only one of 'lemmatize' and 'stemming' can be,
↳ True.")`

`# Validate basic_clean_only option`

```

        if self.basic_clean_only and (lemmatize or lower or remove_stop or
↪remove_punct or remove_num or stemming or
                                add_user_mention_prefix or
↪remove_hashtag_prefix):
            raise ValueError("If 'basic_clean_only' is set to True, other
↪processing options must be set to False.")

        # Assign lemmatize and stemming

        self.lemmatize = lemmatize
        self.stemming = stemming

    def basic_clean(self, text):
        soup = BeautifulSoup(text, "html.parser")
        text = soup.get_text()
        text = re.sub(r'[\n\r]', ' ', text)
        return text.strip()

    def get_cores(self):
        """
        Get the number of CPU cores to use in parallel processing.
        """
        # Get the number of CPU cores available on the system.
        num_cores = os.cpu_count()
        if num_cores < 3:
            use_cores = 1
        else:
            use_cores = num_cores // 2 + 1
        return use_cores

    def spacy_preprocessor(self, texts):
        final_result = []
        nlp = spacy.load(self.model)

        # Disable unnecessary pipelines in spaCy model
        if self.lemmatize:
            # Disable parser and named entity recognition
            disabled_pipes = ['parser', 'ner']
        else:
            # Disable tagger, parser, attribute ruler, lemmatizer and named
↪entity recognition
            disabled_pipes = ['tok2vec', 'tagger', 'parser', 'attribute_ruler',
↪'lemmatizer', 'ner']

        with nlp.select_pipes(disable=disabled_pipes):
            # Modify tokenizer behavior based on user_mention_prefix and
↪hashtag_prefix settings

```

```

        if self.add_user_mention_prefix or self.remove_hashtag_prefix:
            prefixes = list(nlp.Defaults.prefixes)
            if self.add_user_mention_prefix:
                prefixes += ['@'] # Treat '@' as a separate token
            if self.remove_hashtag_prefix:
                prefixes.remove(r'#') # Don't separate '#' from the
→following text
            prefix_regex = spacy.util.compile_prefix_regex(prefixes)
            nlp.tokenizer.prefix_search = prefix_regex.search

        # Process text data in parallel using spaCy's nlp.pipe()
        for doc in nlp.pipe(texts, batch_size=self.batch_size, n_process=self.
→get_cores()):
            filtered_tokens = []
            for token in doc:
                # Check if token should be removed based on specified filters
                if self.remove_stop and token.is_stop:
                    continue
                if self.remove_punct and token.is_punct:
                    continue
                if self.remove_num and token.like_num:
                    continue
                if self.remove_url and token.like_url:
                    continue
                if self.remove_email and token.like_email:
                    continue

                # Append the token's text, lemma, or stemmed form to the
→filtered_tokens list
                if self.lemmatize:
                    filtered_tokens.append(token.lemma_)
                elif self.stemming:
                    filtered_tokens.append(PorterStemmer().stem(token.text))
                else:
                    filtered_tokens.append(token.text)

                # Join the tokens and apply lowercasing if specified
            text = ' '.join(filtered_tokens)
            if self.lower:
                text = text.lower()
            final_result.append(text.strip())

        return final_result

def fit(self, X, y=None):
    return self

```

```

def transform(self, X, y=None):
    try:
        if not isinstance(X, (list, np.ndarray)):
            raise TypeError(f'Expected list or numpy array, got {type(X)}')

        x_clean = [self.basic_clean(text).encode('utf-8', 'ignore').
↪decode() for text in X]

        # Check if only basic cleaning is required
        if self.basic_clean_only:
            return x_clean # Return the list of basic-cleaned texts

        x_clean_final = self.spacy_preprocessor(x_clean)
        return x_clean_final

    except Exception as error:
        print(f'An exception occurred: {repr(error)}')

```

####Featurizer

```

[55]: class ManualFeatures(TransformerMixin, BaseEstimator):

    """A transformer class for extracting manual features from text data.

    This class is designed to be used in a scikit-learn pipeline. It uses the_
↪spaCy
    library to extract a variety of manual features from text data, such as
    part-of-speech (POS) features, named entity recognition (NER) features,
    and count-based features.
    """

    def __init__(self, spacy_model='en_core_web_sm', batch_size = 64,
↪pos_features = True, ner_features = True, count_features = True):

        """
        Initialize the feature extractor.

        Parameters
        -----
        spacy_model : str
            The name of the spaCy model to use for feature extraction.
        pos_features : bool, optional (default=True)
            Whether to extract part-of-speech (POS) features from the text data.
        ner_features : bool, optional (default=True)
        """

```

```

        Whether to extract named entity recognition (NER) features from the
        ↪ text data.
        count_features : bool, optional (default=True)
        Whether to extract count-based features from the text data.
        """

        self.spacy_model = spacy_model
        self.batch_size = batch_size
        self.pos_features = pos_features
        self.ner_features = ner_features
        self.count_features = count_features

    def get_cores(self):
        """
        Get the number of CPU cores to use in parallel processing.
        """
        # Get the number of CPU cores available on the system.
        num_cores = os.cpu_count()
        if num_cores < 3:
            use_cores = 1
        else:
            use_cores = num_cores // 2 + 1
        return num_cores

    def get_pos_features(self, cleaned_text):

        nlp = spacy.load(self.spacy_model)
        noun_count = []
        aux_count = []
        verb_count = []
        adj_count = []

        # Disable the lemmatizer and NER pipelines for improved performance
        disabled_pipes = ['lemmatizer', 'ner']
        with nlp.select_pipes(disable=disabled_pipes):
            n_process = self.get_cores()
            for doc in nlp.pipe(cleaned_text, batch_size=self.batch_size,
            ↪ n_process=n_process):
                # Extract nouns, auxiliaries, verbs, and adjectives from the
                ↪ document
                nouns = [token.text for token in doc if token.pos_ in
                ↪ ["NOUN", "PROPN"]]
                auxs = [token.text for token in doc if token.pos_ in ["AUX"]]
                verbs = [token.text for token in doc if token.pos_ in ["VERB"]]
                adjectives = [token.text for token in doc if token.pos_ in
                ↪ ["ADJ"]]

```

```

        # Store the count of each type of word in separate lists
        noun_count.append(len(nouns))
        aux_count.append(len(auxs))
        verb_count.append(len(verbs))
        adj_count.append(len(adjectives))

    # Stack the count lists vertically to form a 2D numpy array
    return np.transpose(np.vstack((noun_count, aux_count, verb_count,
↪adj_count)))

def get_ner_features(self, cleaned_text):
    nlp = spacy.load(self.spacy_model)
    count_ner = []

    # Disable the tok2vec, tagger, parser, attribute ruler, and lemmatizer,
↪pipelines for improved performance
    disabled_pipes = ['tok2vec', 'tagger', 'parser', 'attribute_ruler',
↪'lemmatizer']
    with nlp.select_pipes(disable=disabled_pipes):
        n_process = self.get_cores()
        for doc in nlp.pipe(cleaned_text, batch_size=self.batch_size,
↪n_process=n_process):
            ners = [ent.label_ for ent in doc.ents]
            count_ner.append(len(ners))

    # Convert the list of NER counts to a 2D numpy array
    return np.array(count_ner).reshape(-1, 1)

def get_count_features(self, cleaned_text):
    list_count_words = []
    list_count_characters = []
    list_count_characters_no_space = []
    list_avg_word_length = []
    list_count_digits = []
    list_count_numbers = []
    list_count_misspell=[]
    list_count_sentences = []

    nlp = spacy.load(self.spacy_model)
    disabled_pipes = ['tok2vec', 'tagger', 'parser', 'attribute_ruler',
↪'lemmatizer', 'ner']
    with nlp.select_pipes(disable=disabled_pipes):
        if not nlp.has_pipe('sentencizer'):
            nlp.add_pipe('sentencizer')

```

```

        n_process = self.get_cores()
        for doc in nlp.pipe(cleaned_text, batch_size=self.batch_size,
↪n_process=n_process):
            count_word = len([token for token in doc if not token.is_punct])
            count_char = len(doc.text)
            count_char_no_space = len(doc.text_with_ws.replace(' ', ''))
            avg_word_length = count_char_no_space / (count_word + 1)
            count_numbers = len([token for token in doc if token.is_digit])
            count_sentences = len(list(doc.sents))

            list_count_words.append(count_word)
            list_count_characters.append(count_char)
            list_count_characters_no_space.append(count_char_no_space)
            list_avg_word_length.append(avg_word_length)
            list_count_numbers.append(count_numbers)
            list_count_sentences.append(count_sentences)

        count_features = np.vstack((list_count_words, list_count_characters,
↪list_count_characters_no_space, list_avg_word_length,
                                list_count_numbers, list_count_sentences))
        return np.transpose(count_features)

    def fit(self, X, y=None):
        """
        Fit the feature extractor to the input data.

        This method does not actually do any fitting, as the feature extractor
↪is stateless.
        It simply returns the instance of the class.

        Parameters:
        X (list or numpy.ndarray): The input data.
        y (list or numpy.ndarray, optional): The target labels. Not used in
↪this implementation.

        Returns:
        FeatureExtractor: The instance of the class.
        """
        return self

    def transform(self, X, y=None):
        """
        Transform the input data into a set of features.

        Parameters:

```



```

    X (list or numpy.ndarray): The input data.
    y (list or numpy.ndarray, optional): The target labels. Not used in
    ↪ this implementation.

    Returns:
        tuple: A tuple containing a 2D numpy array with shape (len(X),
    ↪ num_features) where num_features is the number of features extracted and a
    ↪ list of feature names.

    Raises:
        TypeError: If the input data is not a list or numpy array.
        Exception: If an error occurs while transforming the data into features.
    """
    try:
        # Check if the input data is a list or numpy array
        if not isinstance(X, (list, np.ndarray)):
            raise TypeError(f"Expected list or numpy array, got {type(X)}")

        preprocessor1 = SpacyPreprocessor(model='en_core_web_sm',
    ↪ batch_size=64, lemmatize=False, lower=False,
                                remove_stop=False, remove_email=True,
                                remove_url=True, remove_num=False,
    ↪ stemming=False,
                                add_user_mention_prefix=True,
    ↪ remove_hashtag_prefix=False, basic_clean_only=False)
        preprocessor2 = SpacyPreprocessor(model='en_core_web_sm',
    ↪ batch_size=64, lemmatize=False, lower=False,
                                remove_stop=False, remove_punct=False,
    ↪ remove_email=True,
                                remove_url=True, remove_num=False,
    ↪ stemming=False,
                                add_user_mention_prefix=True,
    ↪ remove_hashtag_prefix=False, basic_clean_only=False)

        feature_names = []
        if self.pos_features or self.ner_features:
            cleaned_x_count_ner_pos = preprocessor2.fit_transform(X)

        if self.count_features:
            cleaned_x_count_features = preprocessor1.fit_transform(X)
            count_features = self.
    ↪ get_count_features(cleaned_x_count_features)
            feature_names.extend(['count_words', 'count_characters',
                                'count_characters_no_space',
    ↪ 'avg_word_length',
                                'count_numbers', 'count_sentences'])

```

```

        else:
            count_features = np.empty(shape=(0, 0))

        if self.pos_features:
            pos_features = self.get_pos_features(cleaned_x_count_ner_pos)
            feature_names.extend(['noun_count', 'aux_count', 'verb_count',
↪ 'adj_count'])
        else:
            pos_features = np.empty(shape=(0, 0))

        if self.ner_features:
            ner_features = self.get_ner_features(cleaned_x_count_ner_pos)
            feature_names.extend(['ner'])
        else:
            ner_features = np.empty(shape=(0, 0))

        # Stack the feature arrays horizontally to form a single 2D numpy
↪ array
        if ner_features.shape == (0, 0) and pos_features.shape == (0, 0):
            return np.hstack((count_features))
        elif pos_features.shape == (0, 0):
            return np.hstack((count_features, ner_features))
        elif ner_features.shape == (0, 0):
            return np.hstack((count_features, pos_features))
        else:
            return np.hstack((count_features, ner_features, pos_features))

    except Exception as error:
        print(f'An exception occured: {repr(error)}')

```

```

[56]: pipe2 = Pipeline([('feature_eng', ManualFeatures()), ('classifier',
↪ XGBClassifier(scale_pos_weight=estimate))]
param_grid = {'feature_eng__pos_features': [True, False],
↪ 'feature_eng__ner_features': [True, False]}

```

```

[57]: random_search2 = RandomizedSearchCV(estimator=pipe2,
↪ param_distributions=param_grid, cv = 2, scoring= f2score)
random_search2.fit(X_train, y_train)

```

```

[57]: RandomizedSearchCV(cv=2,
                        estimator=Pipeline(steps=[('feature_eng', ManualFeatures()),
                                                    ('classifier',
XGBClassifier(base_score=None,
                                                    booster=None,
                                                    callbacks=None,
colsample_bylevel=None,
colsample_bynode=None,

```

```

colsample_bytree=None,
early_stopping_rounds=None,
enable_categorical=False,

min_child_weight=None,

monotone_constraints=None,

num_parallel_tree=None,

...))),
        param_distributions={'feature_eng__ner_features': [True,
                                                             False],
                              'feature_eng__pos_features': [True,
                                                             False]},
        scoring=make_scorer(fbeta_score, beta=2))

```

```

[58]: print("Best Parameters: ", random_search2.best_params_)
      print("Best F2 score:", random_search2.best_score_)
      print("Best Estimator: ", random_search2.best_estimator_)

```

```

Best Parameters: {'feature_eng__pos_features': True,
                  'feature_eng__ner_features': False}
Best F2 score: 0.8621541501976286
Best Estimator: Pipeline(steps=[('feature_eng',
                                  ManualFeatures(ner_features=False)),
                                ('classifier',
                                 XGBClassifier(base_score=None, booster=None, callbacks=None,
                                                colsample_bylevel=None, colsample_bynode=None,
                                                colsample_bytree=None, device=None,
                                                early_stopping_rounds=None,
                                                enable_categorical=False, eval_metric=None,
                                                feature_types=None, gamma=None, grow_policy=None,
                                                importance_type=None,
                                                interaction_constraints=None, learning_rate=None,
                                                max_bin=None, max_cat_threshold=None,
                                                max_cat_to_onehot=None, max_delta_step=None,
                                                max_depth=None, max_leaves=None,

```

```

min_child_weight=None, missing=nan,
monotone_constraints=None, multi_strategy=None,
n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=None,
...)))

```

```

[59]: print("The train score is", random_search2.score(X_train, y_train))
      print("The test score is", random_search2.score(X_test, y_test))

```

```

The train score is 1.0
The test score is 0.85000000000000001

```

4.0.4 Sparse Embeddings + Feature Engineering *italicised text*

```

[60]: feature_extract = FeatureUnion([("sparse_embed", TfidfVectorizer()),
      ↪('feature_eng', ManualFeatures())])

```

```

[61]: pipe3 = Pipeline([('fe', feature_extract), ('classifier',
      ↪XGBClassifier(scale_pos_weight=estimate))])
param_grid = {'fe__sparse_embed__max_df': [0.15, 0.2, 0.5, 0.75, 1],
      ↪('fe__sparse_embed__max_features': [None, 100, 200, 400,
      ↪500], 'fe__feature_eng__pos_features': [True, False],
      ↪('fe__feature_eng__ner_features': [True, False])}

```

```

[62]: random_search3 = RandomizedSearchCV(estimator=pipe3,
      ↪param_distributions=param_grid, cv = 2, scoring= f2score)
random_search3.fit(X_train, y_train)

```

```

[62]: RandomizedSearchCV(cv=2,
      estimator=Pipeline(steps=[('fe',
      FeatureUnion(transformer_list=[('sparse_embed',
      TfidfVectorizer()),
      ('feature_eng',
      ManualFeatures())])),
      ('classifier',
      XGBClassifier(base_score=None,
      booster=None,
      callbacks=None,
      device=None,
      enable_categorical=False,
      multi_strategy=None,
      n_estimators=None,
      n_jobs=None,
      num_parallel_tree=None,
      random_state=None,
      scale_pos_weight=None,
      subsample=None,
      tree_method=None,
      validate_parameters=None,
      verbosity=None,
      watchlist=None)),
      colsample_bylevel=None,
      colsample_bynode=None,
      colsample_bytree=None,
      early_stopping_rounds=None,
      eval_metric=None,
      feature_selector=None,
      ignore_warnings=None,
      keep_training_samples=None,
      learning_rate=None,
      max_delta_step=None,
      max_depth=None,
      max_leaves=None,
      min_child_weight=None,
      missing=None,
      monotone_constraints=None,
      multi_strategy=None,
      n_estimators=None,
      n_jobs=None,
      num_parallel_tree=None,
      random_state=None,
      reg_alpha=None,
      reg_lambda=None,
      sample_weight=None,
      silent=None,
      subsample=None,
      tree_method=None,
      validate_parameters=None,
      verbosity=None,
      watchlist=None))

```

```

num_parallel_tree=None,
                                                    random_state=None,
...))),
        param_distributions={'fe__feature_eng__ner_features': [True,
False],
                            'fe__feature_eng__pos_features': [True,
False],
                            'fe__sparse_embed__max_df': [0.15, 0.2,
0.5, 0.75,
1],
                            'fe__sparse_embed__max_features': [None,
100,
200,
400,
500]}},
        scoring=make_scorer(fbeta_score, beta=2))

```

```

[63]: print("Best Parameters: ", random_search3.best_params_)
      print("Best F2 score:", random_search3.best_score_)
      print("Best Estimator: ", random_search3.best_estimator_)

```

```

Best Parameters: {'fe__sparse_embed__max_features': 100,
'fe__sparse_embed__max_df': 0.75, 'fe__feature_eng__pos_features': True,
'fe__feature_eng__ner_features': True}
Best F2 score: 0.8625428005045954
Best Estimator: Pipeline(steps=[('fe',
                                FeatureUnion(transformer_list=[('sparse_embed',
                                                                TfIdfVectorizer(max_df=0.75,
max_features=100)),
                                                                ('feature_eng',
                                                                ManualFeatures())])),
                                ('classifier',
                                XGBClassifier(base_score=None, booster=None, callbacks=None,
colsample_bylevel=None, colsample_bynode=None,
colsample_bytree=None, device=None,
early_stopping_rounds=None,
enable_categorical=F...
feature_types=None, gamma=None, grow_policy=None,
importance_type=None,
interaction_constraints=None, learning_rate=None,
max_bin=None, max_cat_threshold=None,
max_cat_to_onehot=None, max_delta_step=None,
max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan,
monotone_constraints=None, multi_strategy=None,
n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=None,

```

...)))]

```
[64]: print("The train score is", random_search3.score(X_train, y_train))
      print("The test score is", random_search3.score(X_test, y_test))
```

The train score is 1.0

The test score is 0.9278350515463917

After running the three pipelines using RandomSearchCV, our corresponding validation F2 scores are as follows:

1. Sparse Embeddings -> 0.7359480145431272
2. Feature Engineering -> 0.8621541501976286
3. Sparse Embeddings + Feature Engineering -> 0.8625428005045954

The test scores for the pipelines are as follows:

1. Sparse Embeddings -> 0.7216494845360824
2. Feature Engineering -> 0.8500000000000001
3. Sparse Embeddings + Feature Engineering -> 0.9278350515463917

From these results, we clearly see that the **third pipeline** (Sparse Embeddings + Feature Engineering) has the best test Score as well as the F2 score. We shall thus choose this pipeline to test on the larger dataset.

5 Required Submissions:

1. Submit two colab/jupyter notebooks
 - (analysis with smaller subset and all three pipelines)
 - (analysis with bigger subset and only final pipeline)
2. Pdf version of the notebooks (HWs will not be graded if pdf version is not provided).
3. **The notebooks and pdf files should have the output.**
4. Name files as follows : **FirstName_file1_hw2, FirstName_file2_h2**